

# Service

“ a method for **exposing** a network application that is running as one or more Pods in your cluster”



# Demo

k8s/03\_service/multiple\_services.yaml  
(DNS resolution)



# Service vs Deployment



DevOps perspective:

microservice = k8s Service

Developer perspective:

microservice = k8s Deployment

Service = it's mainly about **networking**

Deployment = it's mainly about a **workload**



# Ingress

“lets you map traffic to **different backends**”

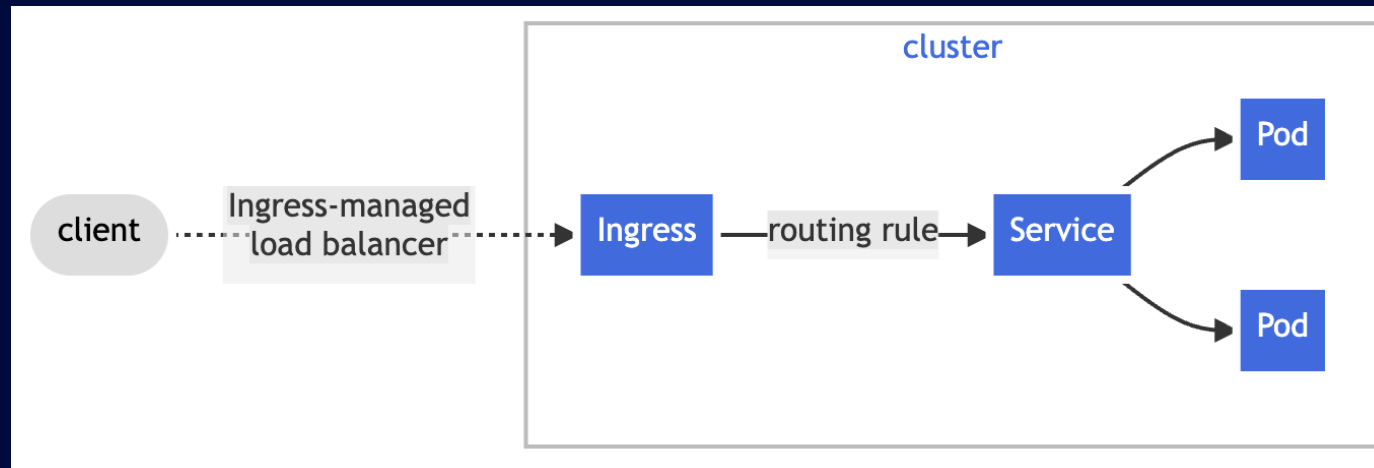
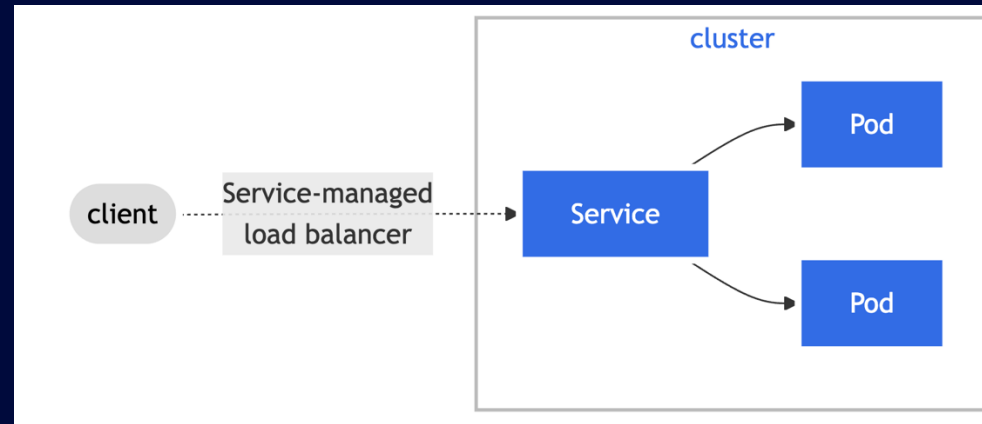


Service = handles just **1 Deployment**

Ingress – handles **1..N Deployments** (via Services)



# Ingress vs Service



# Demo

k8s/04\_ingress/two\_services.yaml





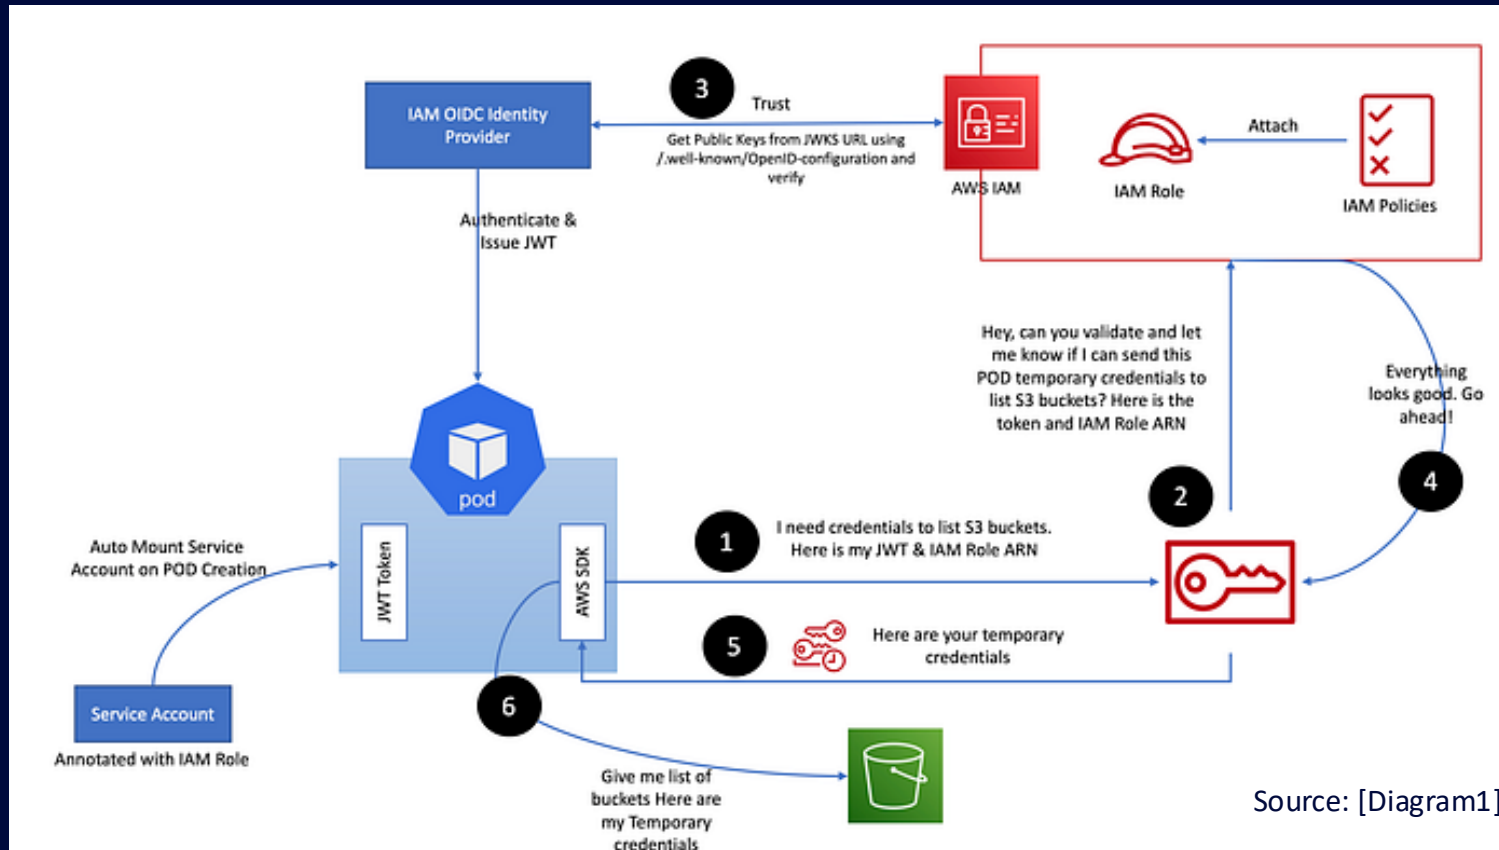
# Service Accounts

“Application Pods, system components, [...] can use a specific ServiceAccount's credentials **to identify as that ServiceAccount.**”

i.e Pods can identify as a specific ServiceAccount



# IRSA = IAM Roles for Service Accounts



# Amazon EKS Pod Identity Webhook

ServiceAccount (sa-test) annotated with:

```
eks.amazonaws.com/role-arn: "arn:aws:iam::{account-id}:role/{iam-role}"
```

(new) Pod using:

```
spec.serviceAccountName: sa-test
```

Injects (by web hook):

Environment variables:

- AWS\_DEFAULT\_REGION
- AWS\_REGION
- AWS\_ROLE\_ARN
- AWS\_WEB\_IDENTITY\_TOKEN\_FILE
- AWS\_STS\_REGIONAL\_ENDPOINTS

Volumes:

- var/run/secrets/eks.amazonaws.com/serviceaccount/token



```
val client = S3Client
    .builder().credentialsProvider(
        StsAssumeRoleWithWebIdentityCredentialsProvider.builder().build()
    )
    .region(US_EAST_1)
    .build()
```



```
private StsWebIdentityTokenFileCredentialsProvider(Builder builder) {
    super(builder, "sts-assume-role-with-web-identity-credentials-provider");
    Path webIdentityTokenFile =
        builder.webIdentityTokenFile != null
            ? builder.webIdentityTokenFile
            : Paths.get(
                trim(
                    SdkSystemSetting.AWS_WEB_IDENTITY_TOKEN_FILE.getStringValueOrThrow()
                )
            );
}
```



# Demo

`aws/create/7_setup_irsa.sh`

`k8s/05_irsa`



# Helm

package manager for Kubernetes



A Chart is a Helm **package**

(Helm uses a **packaging format** called *charts*. A chart is a collection of files that describe a related set of Kubernetes resources.)

[used by UP]

A Repository is the place where charts can be collected and shared

[not used by UP] 

A Release is an **instance of a chart** running in a Kubernetes cluster.

[not used by UP] 



**User Platform facilitates only one command of Helm**  
(under the hood – it is called by ArgoCD):

**helm template**





# GO templates

<https://pkg.go.dev/text/template> (engine)

<https://pkg.go.dev/github.com/Masterminds/sprig> (extra template functions - supported out of the box)

**"Actions"**--data evaluations or control structures--are delimited by **"{{"** and **"}}"**;  
all text **outside actions** is copied to the output **unchanged**.



# Demo

```
aws/create/5_create_ingress_controller.sh  
aws/create/8_install_keda.sh
```

(as regular package manager: `helm install`)



# Demo


k8s/06\_helm  
(helm template)



# ArgoCD

**continuous delivery tool for Kubernetes**





# Demo

k8s/07\_argo



# UP applications setup in EKS

## ApplicationSet – generator #1 (tenants)

(line 23: - path: "applications/eu-west-1/\*.yaml" )

## UP applications - input to #1 generator

## ApplicationSet – generator #2

(line 18: {{- range \$d := .Values.destinations }} )

## UP environments – input to #2 generator



## ApplicationSet (use1-up)







# Demo

apply change in GH repo (**source**) and see  
changes reflected in k8s cluster (**target**)



# UP pull request merge


## utilities/eks-rolling-update-deploy :

1. Cloning repository [https://\[...\]/viacomcbs/up-k8s-applications.git](https://[...]/viacomcbs/up-k8s-applications.git)
2. `python3 -u argocd-update-helm-values.py` 
3. `git commit [...]`
4. `git push --set-upstream origin main`
5. `python3 -u argocd-deployment-validation.py` 

```
def restart_argocd_application(argocd_server, argocd_app, token):  
    token_opt = f"--auth-token={token}"  
    command = " ".join([  
        f"argocd app actions run {argocd_app}",  
        "restart --kind Deployment",  
        f"{token_opt if token != '' else ''}",  
        f"--server {argocd_server}",  
        "--grpc-web",  
    ])
```

```
argocd app actions run use1-up-braze-integration-qa restart \  
  --kind Deployment {token_opt} \  
  --server {argocd_server} \  
  --grpc-web
```

```
{  
    helm_values_file: braze-integration/use1-qa.yaml,  
    key_to_update: pod.containers.braze-integration.image,  
    container_image: [...].[/up-braze-integration:0.1.290,  
    patch_values: {  
        pod: {  
            containers: {  
                braze-integration: {  
                    environmentVariables: {  
                        NR_VERSION: 0.1.290  
                    }  
                }  
            }  
        }  
    }  
}
```

**! Triggers  
Kubernetes  
rolling update** 



# UP service restart (e.g. after configuration change)

user-platform/up-eks-application-restart

utilities/argocd-restart-application

```
argocd app actions run use1-up-braze-integration-qa restart \
--kind Deployment --all '--auth-token=****' \
--server argocd.tools.paramount.tech \
--grpc-web
```


**! Triggers  
Kubernetes  
rolling update**



# UP 'perf' service horizontal scaling

## user-platform/up-scale-perf-eks-services:

1. Cloning repository [https://\[...\]/viacomcbs/up-k8s-applications.git](https://[...]/viacomcbs/up-k8s-applications.git)
2. `sed -i '/replicas:/s/.*:/: 0/' api-gateway/use1-perf.yaml`
3. `git commit [...]`
4. `git push`

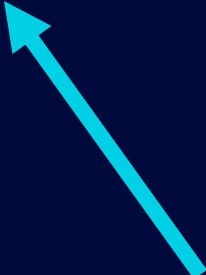


**! ArgoCD autosync  
triggers Kubernetes  
rolling update**



# UP autoscaling (Keda)

1. **Open a PR for the service that needs an autoscaling change:**
  - E.g. <https://github.com/viacomcbs/up-k8s-applications/blob/main/activation-code/use1-prod.yaml>
2. **Merge PR**
3. **Wait for changes to be picked up by ArgoCD autosync**



```
autoscaling:
  keda:
    enabled: true
    maxReplicas: 64
    minReplicas: *replicas
    pollingInterval: 60
    advancedBehavior:
      scaleUp:
        stabilizationWindowSeconds: 120
```



# Sources

Docs1 - <https://kubernetes.io/docs/home/>

Docs2 - <https://helm.sh/docs/topics/charts/>

Docs3 - <https://argo-cd.readthedocs.io/en/stable/>

Diagram1 [IRSA] - <https://mohaamer5.medium.com/iam-roles-for-service-accounts-with-eks-irsa-good-bye-aws-credentials-1cdf1fa5192>



# Terraform

